

#11

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import PolynomialFeatures
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.metrics import mean_squared_error, r2_score
```

```
# Step 1: Generate a synthetic dataset
```

```
np.random.seed(0)
```

```
X = np.random.rand(100, 1) * 10 # Feature matrix
```

```
y = 2 * (X ** 2) + 3 * X + 5 + np.random.randn(100, 1) * 5 # Target vector
```

```
# Step 2: Split the dataset into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
# Step 3: Preprocess the data to include polynomial features
```

```
poly = PolynomialFeatures(degree=2) # Try different degrees for different results
```

```
X_poly_train = poly.fit_transform(X_train)
```

```
X_poly_test = poly.transform(X_test)
```

```
# Step 4: Fit a polynomial regression model
```

```
model = LinearRegression()
```

```
model.fit(X_poly_train, y_train)
```

```
# Step 5: Predict and evaluate the model
```

```
y_train_pred = model.predict(X_poly_train)
```

```
y_test_pred = model.predict(X_poly_test)
```

```
# Performance metrics
```

```
train_mse = mean_squared_error(y_train, y_train_pred)
```

```
test_mse = mean_squared_error(y_test, y_test_pred)
```

```
train_r2 = r2_score(y_train, y_train_pred)
```

```
test_r2 = r2_score(y_test, y_test_pred)
```

```
print(f"Train MSE: {train_mse}")
```

```
print(f"Test MSE: {test_mse}")
```

```
print(f"Train R2 Score: {train_r2}")
```

```
print(f"Test R2 Score: {test_r2}")
```

```
# Step 6: Visualize the results
```

```
plt.scatter(X, y, color='blue', label='Actual data')
```

```
X_range = np.linspace(min(X), max(X), 100).reshape(-1, 1)
```

```
y_range_pred = model.predict(poly.transform(X_range))
```

```
plt.plot(X_range, y_range_pred, color='red', label='Polynomial regression fit')
```

```
plt.xlabel('X')
```

```
plt.ylabel('y')
```

```
plt.title('Polynomial Regression')
```

```
plt.legend()
```

```
plt.show()
```


```
#output
```

```
Train MSE: 24.125995439736805
```

```
Test MSE: 25.7136695038184
```

```
Train R2 Score: 0.9942709761929696
```

```
Test R2 Score: 0.9944791639465156
```



#12

```
import numpy as np
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
# Step 1: Load the dataset
```

```
from sklearn.datasets import load_iris
```

```
iris = load_iris()
```

```
X = iris.data
```

```
y = iris.target
```

```
# Step 2: Preprocess the data (if necessary)
```

```
# In this case, the Iris dataset is already clean and ready to use.
```

```
# Step 3: Split the dataset into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Step 4: Fit a KNN model to the training data
```

```
k = 3 # Number of neighbors
```

```
knn = KNeighborsClassifier(n_neighbors=k)
```

```
knn.fit(X_train, y_train)
```

```
# Step 5: Predict the labels for the test data
```

```
y_pred = knn.predict(X_test)
```

```
# Step 6: Evaluate the model's performance
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Accuracy: {accuracy}")
```

```
print("\nClassification Report:")
```

```
print(classification_report(y_test, y_pred))
```

```
# Confusion matrix
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=iris.target_names,  
yticklabels=iris.target_names)
```

```
plt.xlabel('Predicted')
```

```
plt.ylabel('True')
```

```
plt.title('Confusion Matrix')
```


```
plt.show()
```

```
#output
```

```
Accuracy: 1.0
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30



#13

```
import pandas as pd
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.naive_bayes import GaussianNB
```

```
from sklearn.metrics import accuracy_score
```

```
# Step 1: Load the dataset
```

```
data = {
```

```
    'Occupation': ['Engineer', 'Doctor', 'Artist', 'Engineer', 'Artist', 'Doctor', 'Engineer', 'Artist', 'Doctor',  
                  'Engineer'],
```

```
    'Credit Score': [720, 680, 650, 700, 710, 690, 730, 640, 660, 750]
```

```
}
```

```
df = pd.DataFrame(data)
```

```
# a. Print the first five rows
```

```
print("First five rows of the dataset:")
```

```
print(df.head())
```

```
# b. Basic statistical computations
```

```
print("\nBasic statistical computations:")
```

```
print(df.describe())
```

```
# c. The columns and their data types
```

```
print("\nColumns and their data types:")
```

```
print(df.dtypes)
```

```
# d. Detect and handle null values
```

```
print("\nNull values in the dataset:")
```

```
print(df.isnull().sum())
```

As an example, let's manually insert a null value and then handle it

```
df.at[2, 'Credit Score'] = None
```

```
print("\nNull values after insertion:")
```

```
print(df.isnull().sum())
```

Replace null values with the mode

```
mode_value = df['Credit Score'].mode()[0]
```

```
df['Credit Score'].fillna(mode_value, inplace=True)
```

```
print("\nDataset after handling null values:")
```

```
print(df)
```

e. Explore the dataset using a box plot

```
sns.boxplot(x='Occupation', y='Credit Score', data=df)
```

```
plt.title('Credit Scores Based on Occupation')
```

```
plt.show()
```

f. Split the dataset into train and test sets

```
X = pd.get_dummies(df['Occupation'], drop_first=True) # One-hot encoding for categorical variable
```

```
y = df['Credit Score']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

g. Fit the Naive Bayes classifier

```
nb = GaussianNB()
```

```
nb.fit(X_train, y_train)
```

i. Predict the model

```
y_pred = nb.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print("\nPredicted values for the test set:")
```

```
print(y_pred)

print(f"Model Accuracy: {accuracy}")
```

#output

First five rows of the dataset:

	Occupation	Credit Score
0	Engineer	720
1	Doctor	680
2	Artist	650
3	Engineer	700
4	Artist	710

Basic statistical computations:

	Credit Score
count	10.00000
mean	693.00000
std	35.91657
min	640.00000
25%	665.00000
50%	695.00000
75%	717.50000
max	750.00000

Columns and their data types:

Occupation	object
Credit Score	int64

dtype: object

Null values in the dataset:

Occupation	0
Credit Score	0

dtype: int64

Null values after insertion:


Occupation 0

Credit Score 1

dtype: int64

Dataset after handling null values:

	Occupation	Credit Score
0	Engineer	720.0
1	Doctor	680.0
2	Artist	640.0
3	Engineer	700.0
4	Artist	710.0
5	Doctor	690.0
6	Engineer	730.0
7	Artist	640.0
8	Doctor	660.0
9	Engineer	750.0



#14

```
import pandas as pd
```

```
import numpy as np
```

```
# Sample training data as a list of dictionaries
```

```
data = [
```

```
    {'Origin': 'Japan', 'Manufacturer': 'Honda', 'Color': 'Blue', 'Decade': '1980', 'Type': 'Economy',  
    'Example Type': 'Positive'},
```

```
    {'Origin': 'Japan', 'Manufacturer': 'Toyota', 'Color': 'Green', 'Decade': '1970', 'Type': 'Sports',  
    'Example Type': 'Negative'},
```

```
    {'Origin': 'Japan', 'Manufacturer': 'Toyota', 'Color': 'Blue', 'Decade': '1990', 'Type': 'Economy',  
    'Example Type': 'Positive'},
```

```
    {'Origin': 'USA', 'Manufacturer': 'Chrysler', 'Color': 'Red', 'Decade': '1980', 'Type': 'Economy',  
    'Example Type': 'Negative'},
```

```
    {'Origin': 'Japan', 'Manufacturer': 'Honda', 'Color': 'White', 'Decade': '1980', 'Type': 'Economy',  
    'Example Type': 'Positive'}  
]
```

```

# Convert the data into a DataFrame
df = pd.DataFrame(data)

# Extract attributes and target
attributes = df.columns[:-1] # All columns except the last one
target = df.columns[-1] # The last column

# Initialize the specific and general hypotheses
S = ['0'] * len(attributes)
G = [['?'] * len(attributes)]

# Candidate Elimination algorithm
for i, row in df.iterrows():
    if row[target] == 'Positive':
        for j in range(len(attributes)):
            if S[j] == '0':
                S[j] = row[j]
            elif S[j] != row[j]:
                S[j] = '?'
        G = [g for g in G if all((g[j] == '?' or g[j] == row[j]) for j in range(len(attributes)))]
    else:
        new_G = []
        for g in G:
            for j in range(len(attributes)):
                if g[j] == '?':
                    for value in df[attributes[j]].unique():
                        if value != row[j]:
                            new_g = g.copy()
                            new_g[j] = value
                            if all((S[k] == '?' or new_g[k] == '?' or S[k] == new_g[k]) for k in range(len(attributes))):

```

```

        new_G.append(new_g)

G = new_G

# Output the final S and G
print("Final specific hypothesis (S):", S)
print("Final general hypotheses (G):", G)

#output
Final specific hypothesis (S): ['Japan', '?', '?', '?', 'Economy']
Final general hypotheses (G): [['Japan', '?', '?', '?', 'Economy'], ['?', 'Honda', '?', '?', 'Economy']]

#15
import numpy as np
import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score

# Generate synthetic data
np.random.seed(0)
X = 2 - 3 * np.random.normal(0, 1, 100)
y = X - 2 * (X ** 2) + np.random.normal(-3, 3, 100)

X = X[:, np.newaxis]

# Split the data into training/testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)

# 1. Linear Regression
linear_regressor = LinearRegression()
linear_regressor.fit(X_train, y_train) # Train the model

```

```
# Predict using the linear model
y_pred_linear = linear_regressor.predict(X_test)

# 2. Polynomial Regression (degree 2)
polynomial_features = PolynomialFeatures(degree=2)
X_poly = polynomial_features.fit_transform(X_train)

poly_regressor = LinearRegression()
poly_regressor.fit(X_poly, y_train) # Train the model

# Predict using the polynomial model
X_test_poly = polynomial_features.transform(X_test)
y_pred_poly = poly_regressor.predict(X_test_poly)

# Evaluate the models
mse_linear = mean_squared_error(y_test, y_pred_linear)
r2_linear = r2_score(y_test, y_pred_linear)

mse_poly = mean_squared_error(y_test, y_pred_poly)
r2_poly = r2_score(y_test, y_pred_poly)

print(f"Linear Regression MSE: {mse_linear}, R2: {r2_linear}")
print(f"Polynomial Regression MSE: {mse_poly}, R2: {r2_poly}")

# Visualize the results
plt.figure(figsize=(14, 5))

# Plot Linear Regression results
plt.subplot(1, 2, 1)
plt.scatter(X_test, y_test, color='black', label='Data')
```

```
plt.plot(X_test, y_pred_linear, color='blue', linewidth=2, label='Linear Regression')  
plt.xlabel('X')  
plt.ylabel('y')  
plt.title('Linear Regression')  
plt.legend()
```

```
# Plot Polynomial Regression results
```

```
plt.subplot(1, 2, 2)  
plt.scatter(X_test, y_test, color='black', label='Data')  
# Sort the values of X_test for a smoother curve  
sorted_X_test = np.sort(X_test, axis=0)  
sorted_y_pred_poly = poly_regressor.predict(polynomial_features.transform(sorted_X_test))  
plt.plot(sorted_X_test, sorted_y_pred_poly, color='red', linewidth=2, label='Polynomial Regression')  
plt.xlabel('X')  
plt.ylabel('y')  
plt.title('Polynomial Regression (degree=2)')  
plt.legend()  
  
plt.show()
```

```
#output
```

```
Linear Regression MSE: 613.3502296050843, R2: 0.2916278819291981
```

```
Polynomial Regression MSE: 8.695598709479661, R2: 0.9899572554498042
```

