

#16

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from scipy.stats import multivariate_normal
```

```
# Generate synthetic data
```

```
np.random.seed(42)
```

```
n_samples = 300
```

```
mean1 = [0, 0]
```

```
cov1 = [[1, 0.1], [0.1, 1]]
```

```
mean2 = [5, 5]
```

```
cov2 = [[1, -0.1], [-0.1, 1]]
```

```
X = np.vstack([
```

```
    np.random.multivariate_normal(mean1, cov1, n_samples),
```

```
    np.random.multivariate_normal(mean2, cov2, n_samples)
```

```
])
```

```
# Number of components
```

```
k = 2
```

```
# Initialize the parameters
```

```
np.random.seed(42)
```

```
pi = np.ones(k) / k # Mixing coefficients
```

```
means = np.random.rand(k, 2) # Means of the Gaussians
```

```
covariances = np.array([np.eye(2)] * k) # Covariances of the Gaussians
```

```
def e_step(X, pi, means, covariances):
```

```
    N = X.shape[0]
```

```
    r = np.zeros((N, k))
```

```

for i in range(k):
    r[:, i] = pi[i] * multivariate_normal.pdf(X, mean=means[i], cov=covariances[i])
r = r / r.sum(axis=1, keepdims=True)
return r

```

```

def m_step(X, r):
    N, D = X.shape
    pi = r.sum(axis=0) / N
    means = np.dot(r.T, X) / r.sum(axis=0)[:, np.newaxis]
    covariances = np.zeros((k, D, D))

    for i in range(k):
        diff = X - means[i]
        covariances[i] = np.dot(r[:, i] * diff.T, diff) / r[:, i].sum()

    return pi, means, covariances

```

```

def log_likelihood(X, pi, means, covariances):
    N = X.shape[0]
    log_likelihood = 0

    for i in range(k):
        log_likelihood += pi[i] * multivariate_normal.pdf(X, mean=means[i], cov=covariances[i])
    return np.log(log_likelihood).sum()

```

Run the EM algorithm

```
max_iter = 100
```

```
tol = 1e-6
```

```
log_likelihoods = []
```

```
for iteration in range(max_iter):
```

```

r = e_step(X, pi, means, covariances)
pi, means, covariances = m_step(X, r)
log_likelihoods.append(log_likelihood(X, pi, means, covariances))

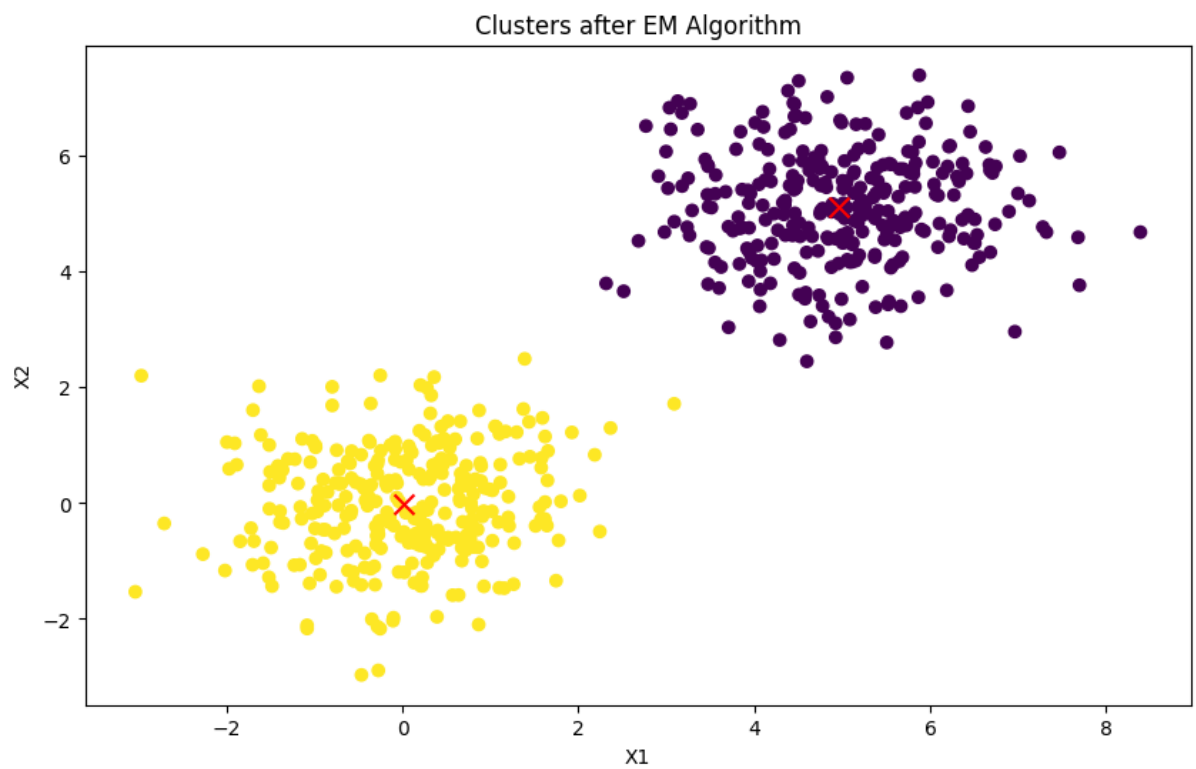
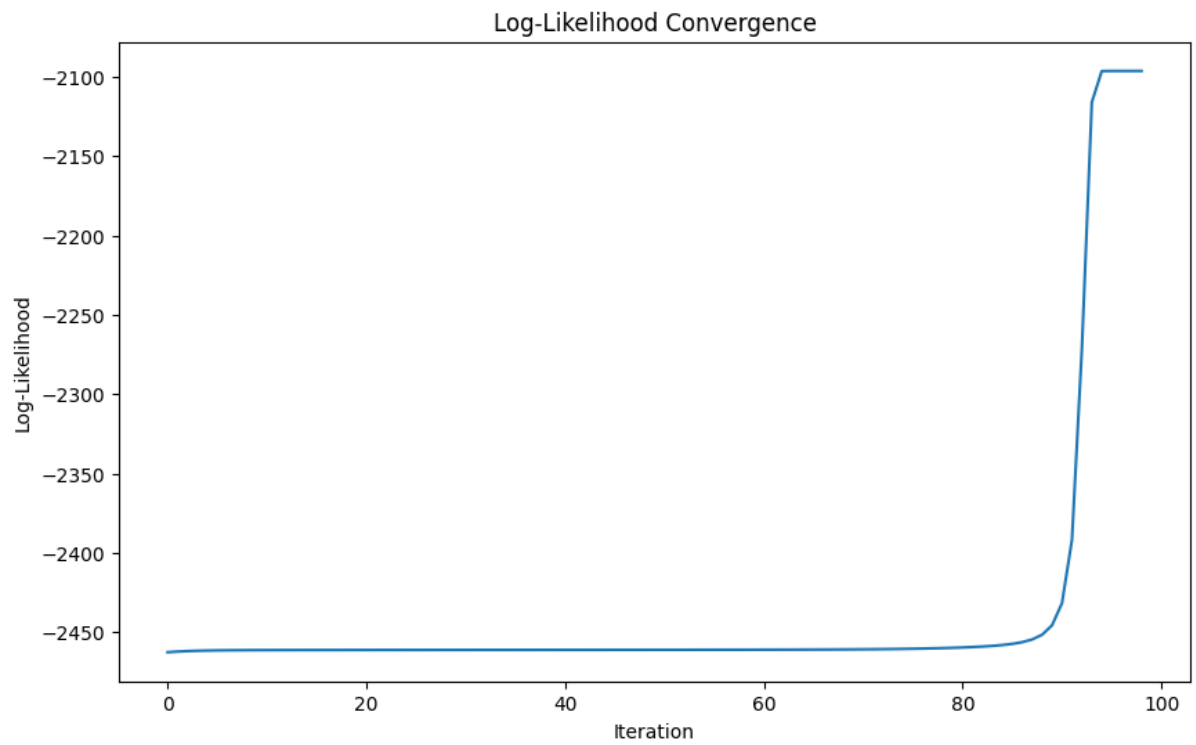
# Check for convergence
if iteration > 0 and abs(log_likelihoods[-1] - log_likelihoods[-2]) < tol:
    break

# Plot the log-likelihoods
plt.figure(figsize=(10, 6))
plt.plot(log_likelihoods)
plt.xlabel('Iteration')
plt.ylabel('Log-Likelihood')
plt.title('Log-Likelihood Convergence')
plt.show()

# Plot the final clusters
plt.figure(figsize=(10, 6))
plt.scatter(X[:, 0], X[:, 1], c=r.argmax(axis=1), cmap='viridis', marker='o')
plt.scatter(means[:, 0], means[:, 1], c='red', marker='x', s=100)
plt.title('Clusters after EM Algorithm')
plt.xlabel('X1')
plt.ylabel('X2')
plt.show()

#output

```



#17

```
import pandas as pd
```

```
import numpy as np
```

```

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score


from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
df['species'] = iris.target
df['species'] = df['species'].map({0: 'setosa', 1: 'versicolor', 2: 'virginica'})


# Step 2: Plot the data using a scatter plot "sepal_width" versus "sepal_length" and color species
plt.figure(figsize=(10, 6))
colors = {'setosa': 'red', 'versicolor': 'green', 'virginica': 'blue'}
plt.scatter(df['sepal width (cm)'], df['sepal length (cm)'], c=df['species'].apply(lambda x: colors[x]),
            label=colors)
plt.xlabel('Sepal Width (cm)')
plt.ylabel('Sepal Length (cm)')
plt.title('Sepal Width vs Sepal Length')
plt.legend(colors)
plt.show()


# Step 3: Split the data
X = df.drop(columns='species')
y = df['species']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Step 4: Fit the data to the model
model = LogisticRegression(max_iter=200)
model.fit(X_train, y_train)

```

```

# Evaluate the model

y_pred_train = model.predict(X_train)

y_pred_test = model.predict(X_test)

print(f'Training Accuracy: {accuracy_score(y_train, y_pred_train):.2f}')

print(f'Testing Accuracy: {accuracy_score(y_test, y_pred_test):.2f}')


# Step 5: Predict the model with new test data [5, 3, 1, 0.3]

new_sample = np.array([[5, 3, 1, 0.3]])


prediction = model.predict(new_sample)

predicted_species = prediction[0]

print(f'The predicted species for the new sample [5, 3, 1, 0.3] is: {predicted_species}')

```

#output



Training Accuracy: 0.97

Testing Accuracy: 1.00

The predicted species for the new sample [5, 3, 1, 0.3] is: setosa

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names

```
warnings.warn(
```

```
#18
```

```
import pandas as pd
```

```
# Load the CSV file
```

```
data = {
```

```
    'Shape':['Circular','Circular','Oval','Oval'],
```

```
    'Size' :['Large','Large','Large','Large'],
```

```
    'Color':['Light','Light','Dark','Light'],
```

```
    'Surface' :['Smooth','Irregular','Smooth','Irregular'],
```

```
    'Thickness' :['Thick','Thick','Thin','Thick'],
```

```
    'Target Concept' :['Maligant(+)','Malignant(+)','Benign(-)','Malignant(+)']
```

```
}
```

```
df = pd.DataFrame(data)
```

```
# Initialize S and G
```

```
# S is initialized to the most specific hypothesis
```

```
# G is initialized to the most general hypothesis
```

```
S = ['0'] * (len(df.columns) - 2)
```

```
G = [['?' for _ in range(len(df.columns) - 2)]]
```

```
# Function to update S
```

```
def update_S(s, example):
```

```
    for i in range(len(s)):
```

```
        if s[i] == '0':
```

```
            s[i] = example[i]
```

```
        elif s[i] != example[i]:
```

```
            s[i] = '?'
```

```
    return s
```

```
# Function to update G
```

```
def update_G(g, s, example):
```

```
    new_g = []
```

```
    for h in g:
```

```
        consistent = True
```

```
        for i in range(len(h)):
```

```
            if h[i] != '?' and h[i] != example[i]:
```

```
                consistent = False
```

```
                break
```

```
    if not consistent:
```

```
        for i in range(len(h)):
```

```
            if h[i] == '?':
```

```
                new_h = h[:]
```

```
                new_h[i] = s[i]
```

```
                if new_h not in new_g:
```

```
                    new_g.append(new_h)
```

```
    return new_g
```

```
# Process each training example
```

```
for index, row in df.iterrows():
```

```
    example = list(row[:-1])
```

```
    target = row[-1]
```

```
    if target == 'Malignant (+)':
```

```
        # Positive example
```

```
        S = update_S(S, example)
```

```
        G = [h for h in G if all(h[i] == '?' or h[i] == example[i] for i in range(len(h)))]
```

```
    else:
```

```
        # Negative example
```

```
        G = update_G(G, S, example)
```

```
        S = [s for s in S if not all(s[i] == '?' or s[i] == example[i] for i in range(len(s)))]
```



```
# Output the final S and G
print("Final specific hypothesis S:", S)
print("Final general hypotheses G:", G)
```

```
#output
```

```
Final specific hypothesis S: ['0', '0', '0', '0']
```

```
Final general hypotheses G: []
```

```
#19
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.metrics import mean_squared_error
```

```
# Step 1: Generate synthetic data
```

```
np.random.seed(0)
```

```
X = 2 * np.random.rand(100, 1)
```

```
y = 3 + 4 * X + np.random.randn(100, 1)
```

```
# Step 2: Implement Linear Regression
```

```
def linear_regression(X, y):
```

```
    # Add bias term
```

```
    X_b = np.c_[np.ones((len(X), 1)), X]
```

```
    # Normal equation:  $\theta = (X^T X)^{-1} X^T y$ 
```

```
    theta_best = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(y)
```

```
    return theta_best
```

```
# Fit linear regression model
```

```
theta_best = linear_regression(X, y)
```

```
# Extract coefficients
```

```
intercept, slope = theta_best[0], theta_best[1]
```

```
# Step 3: Show Performance

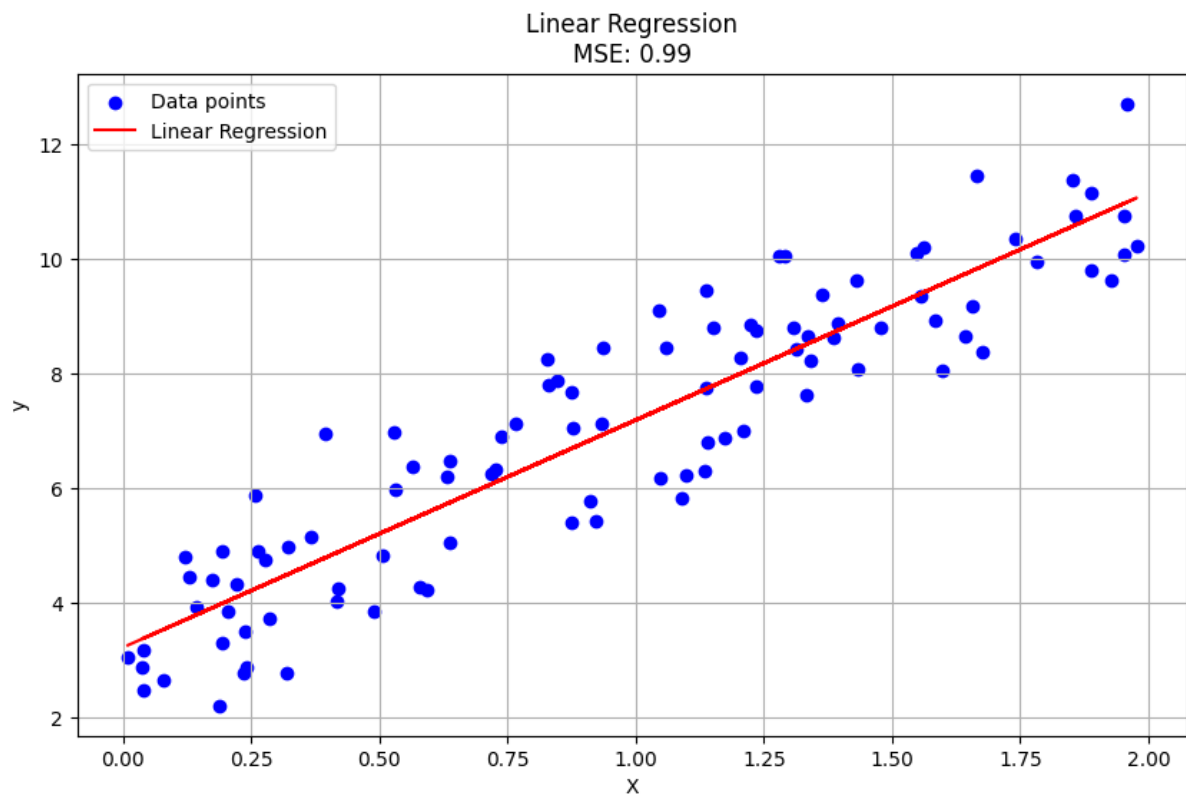
# Predictions
y_predict = np.dot(np.c_[np.ones((len(X), 1)), X], theta_best)

# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y, y_predict)

# Plotting
plt.figure(figsize=(10, 6))
plt.scatter(X, y, color='blue', label='Data points')
plt.plot(X, y_predict, color='red', label='Linear Regression')
plt.title(f'Linear Regression\nMSE: {mse:.2f}')
plt.xlabel('X')
plt.ylabel('y')
plt.legend()
plt.grid(True)
plt.show()

print(f'Intercept (theta_0): {intercept[0]:.2f}')
print(f'Slope (theta_1): {slope[0]:.2f}')
print(f'Mean Squared Error (MSE): {mse:.2f}')

#output
```



Intercept (theta_0): 3.22

Slope (theta_1): 3.97

Mean Squared Error (MSE): 0.99

#20

```
import numpy as np
```

```
from collections import defaultdict
```

```
class NaiveBayes:
```

```
    def __init__(self):
```

```
        self.classes = None
```

```
        self.class_priors = None
```

```
        self.feature_probs = None
```

```
    def fit(self, X, y):
```

```
        self.classes = np.unique(y)
```

```
        self.class_priors = np.bincount(y) / len(y) # Class probabilities (prior)
```

```
self.feature_probs = {c: defaultdict(lambda: np.zeros(X.shape[1])) for c in self.classes} # Initialize as np.zeros
```

```
for c in self.classes:
```

```
    X_class = X[y == c]
```

```
    for i in range(X.shape[1]): # Iterate over features
```

```
        self.feature_probs[c][i] += np.sum(X_class[:, i]) # Update feature counts
```

```
    for i in range(X.shape[1]):
```

```
        self.feature_probs[c][i] /= (X_class.shape[0] + X.shape[1]) # Apply Laplace smoothing
```

```
def predict(self, X):
```

```
    if self.classes is None:
```

```
        raise Exception("Model not fitted yet. Call fit(X, y) first.")
```

```
    y_pred = []
```

```
    for x in X:
```

```
        posteriors = {}
```

```
        for c in self.classes:
```

```
            prior = self.class_priors[c]
```

```
            likelihood = np.prod([self.feature_probs[c][i]**x[i] * (1 - self.feature_probs[c][i])** (1 - x[i])  
for i in range(len(x))])
```

```
            posteriors[c] = prior * likelihood
```

```
        y_pred.append(max(posteriors, key=posteriors.get))
```

```
    return np.array(y_pred)
```

```
# Example usage
```

```
# Sample email data (presence/absence of words) and labels
```

```
X = np.array([
```

```
    [1, 1, 0, 0, 0, 0], # "free money" email (spam)
```

```
    [1, 0, 0, 1, 1, 0], # Meeting invitation email (not spam)
```

```
    [0, 0, 1, 0, 0, 1], # Project update email (not spam)
```

```

    [1, 1, 0, 1, 0, 0], # "Get rich quick" email (spam)
])

y = np.array([1, 0, 0, 1]) # 1 for spam, 0 for not spam

# Create and train the Naive Bayes model
model = NaiveBayes()
model.fit(X, y)

# New email to classify (presence/absence of words)
new_email = np.array([0, 1, 0, 1, 1, 1]) # Email about a meeting and deadline

# Predict class label for the new email
predicted_class = model.predict(new_email.reshape(1, -1)) # Reshape for single data point
print("Predicted class:", "spam" if predicted_class[0] == 1 else "not spam")

#output
Predicted class: not spam

#21
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

# Sample data
data = {
    'size': [1400, 1600, 1700, 1800, 1900, 2000],

```

```
'bedrooms': [3, 3, 4, 4, 4, 5],  
'bathrooms': [2, 3, 2, 3, 3, 4],  
'location': [1, 1, 2, 2, 3, 3],  
'price': [300000, 350000, 400000, 420000, 450000, 500000]  
}
```

```
# Convert to DataFrame
```

```
df = pd.DataFrame(data)
```

```
# Step 1: Print the first five rows of the dataset
```

```
print("First five rows of the dataset:")
```

```
print(df.head())
```

```
# Step 2: Basic statistical computations
```

```
print("\nBasic statistical computations:")
```

```
print(df.describe())
```

```
# Step 3: Print columns and their data types
```

```
print("\nColumns and their data types:")
```

```
print(df.dtypes)
```

```
# Step 4: Detect and handle null values
```

```
print("\nDetecting null values:")
```

```
print(df.isnull().sum())
```

```
print("\nNull values after replacement:")
```

```
print(df.isnull().sum())
```

```
plt.figure(figsize=(10, 8))
```

```
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
```

```
plt.title('Heatmap of Feature Correlations')
```

```
plt.show()
```

```
X = df.drop('price', axis=1)
```

```
y = df['price']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
rmse = np.sqrt(mse)
```

```
print("\nModel Performance:")
```

```
print(f"Root Mean Squared Error: {rmse}")
```

```
new_house = pd.DataFrame({
```

```
    'size': [1500],
```

```
    'bedrooms': [3],
```

```
    'bathrooms': [2],
```

```
    'location': [1]
```

```
})
```

```
predicted_price = model.predict(new_house)
```

```
print("\nPredicted price for the new house:", predicted_price[0])
```

```
#output
```

First five rows of the dataset:

```
size bedrooms bathrooms location price
```

0	1400	3	2	1	300000
1	1600	3	3	1	350000
2	1700	4	2	2	400000
3	1800	4	3	2	420000
4	1900	4	3	3	450000

Basic statistical computations:

	size	bedrooms	bathrooms	location	price
count	6.000000	6.000000	6.000000	6.000000	6.000000
mean	1733.333333	3.833333	2.833333	2.000000	403333.333333
std	216.024690	0.752773	0.752773	0.894427	71180.521680
min	1400.000000	3.000000	2.000000	1.000000	300000.000000
25%	1625.000000	3.250000	2.250000	1.250000	362500.000000
50%	1750.000000	4.000000	3.000000	2.000000	410000.000000
75%	1875.000000	4.000000	3.000000	2.750000	442500.000000
max	2000.000000	5.000000	4.000000	3.000000	500000.000000

Columns and their data types:

```
size      int64
bedrooms  int64
bathrooms int64
location  int64
price     int64
dtype: object
```

Detecting null values:

```
size      0
bedrooms  0
bathrooms 0
location  0
price     0
```


dtype: int64

Null values after replacement:

size 0

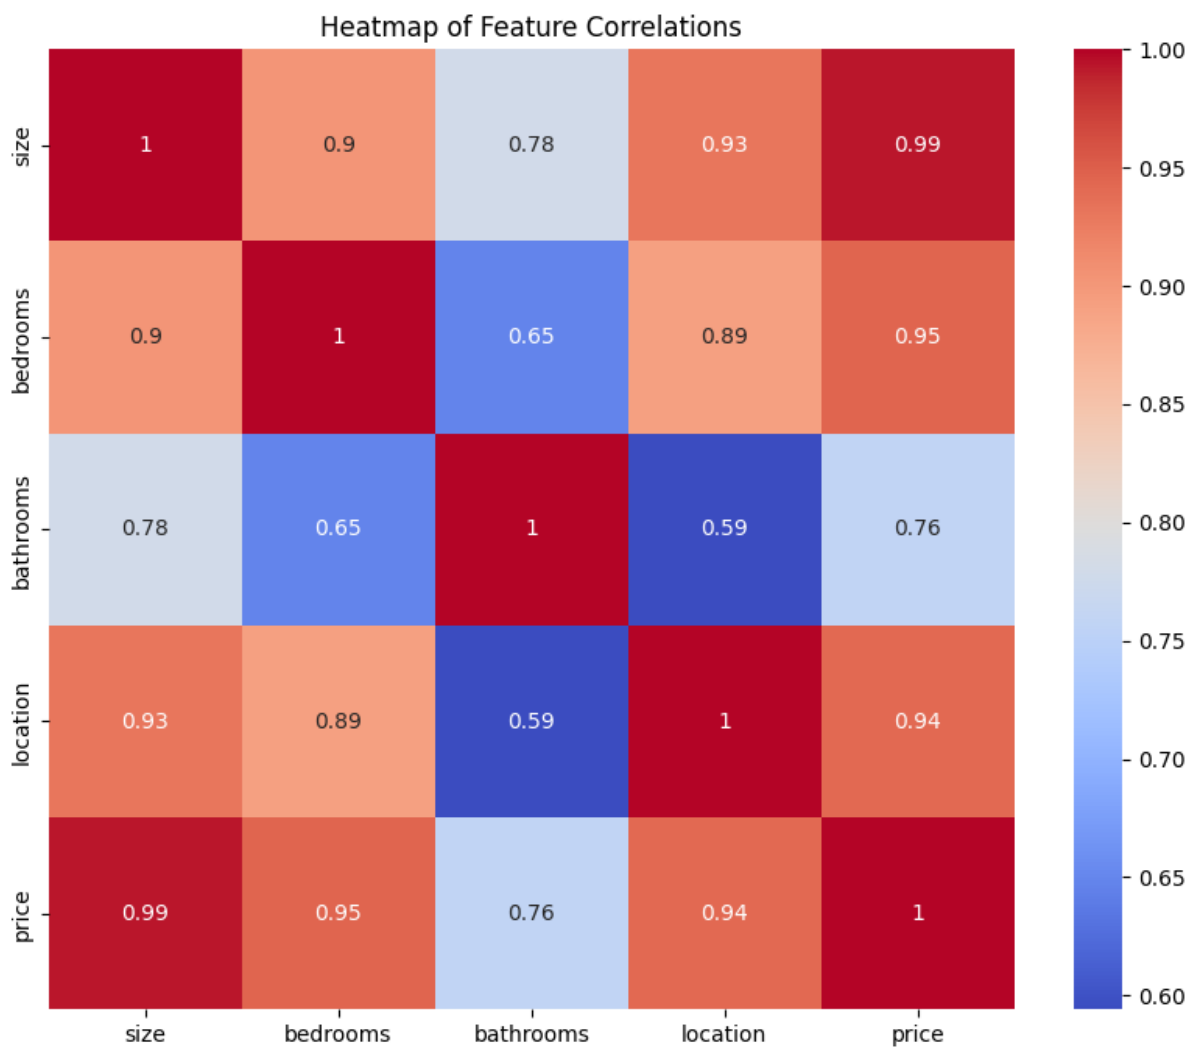
bedrooms 0

bathrooms 0

location 0

price 0

dtype: int64



Model Performance:

Root Mean Squared Error: 12745.832891312073

Predicted price for the new house: 315001.249937503

#22

```
import pandas as pd
```

```
# Function to check if a hypothesis is consistent with an example
```

```
def is_consistent(hypothesis, example):
```

```
    for h, e in zip(hypothesis, example):
```

```
        if h != '?' and h != e:
```

```
            return False
```

```
    return True
```

```
# Function to find the minimal generalization of S
```

```
def generalize_minimally(h, x):
```

```
    new_h = list(h)
```

```
    for i in range(len(h)):
```

```
        if not is_consistent([h[i]], [x[i]]):
```

```
            new_h[i] = '?' if h[i] != x[i] else x[i]
```

```
    return tuple(new_h)
```

```
# Function to find the minimal specialization of G
```

```
def specialize_minimally(h, domains, x):
```

```
    specializations = []
```

```
    for i in range(len(h)):
```

```
        if h[i] == '?':
```

```
            for val in domains[i]:
```

```
                if val != x[i]:
```

```
                    new_h = list(h)
```

```
                    new_h[i] = val
```

```
                    specializations.append(tuple(new_h))
```

```
    elif h[i] != x[i]:
```

```

    new_h = list(h)
    new_h[i] = '?'
    specializations.append(tuple(new_h))
return specializations

```

Candidate-Elimination Algorithm

```
def candidate_elimination(examples):
```

```
    domains = [set(examples[col]) for col in examples.columns[:-1]]
```

```
    n_features = len(domains)
```

Initialize S to the most specific hypothesis

```
S = tuple(['∅'] * n_features)
```

Initialize G to the most general hypothesis

```
G = [tuple(['?'] * n_features)]
```

```
for index, row in examples.iterrows():
```

```
    x, y = row[:-1], row[-1]
```

```
    x = tuple(x)
```

```
    if y == 'yes': # Positive example
```

```
        # Remove from G any hypothesis inconsistent with x
```

```
        G = [g for g in G if is_consistent(g, x)]
```

```
        S = [s for s in S if is_consistent(s, x)]
```

```
        for s in S:
```

```
            if not is_consistent(s, x):
```

```
                S.remove(s)
```

```
                S.append(generalize_minimally(s, x))
```

```
        # Remove from S any hypothesis that is more general than another hypothesis in S
```

```
S = [s for s in S if not any(s != s2 and is_consistent(s2, s) for s2 in S)]
```

```
else: # Negative example
```

```
# Remove from S any hypothesis inconsistent with x
```

```
S = [s for s in S if not is_consistent(s, x)]
```

```
# For each hypothesis g in G that is consistent with x, remove g from G
```

```
# Add to G all minimal specializations h of g such that h is not consistent with x and some  
member of S is more specific than h
```

```
new_G = []
```

```
for g in G:
```

```
    if is_consistent(g, x):
```

```
        new_G.extend(specialize_minimally(g, domains, x))
```

```
    else:
```

```
        new_G.append(g)
```

```
G = new_G
```

```
# Remove from G any hypothesis that is more specific than another hypothesis in G
```

```
G = [g for g in G if not any(g != g2 and is_consistent(g, g2) for g2 in G)]
```

```
return S, G
```

```
data = {
```

```
    'Sky': ['Sunny', 'Sunny', 'Rainy', 'Sunny', 'Sunny'],
```

```
    'AirTemp': ['Warm', 'Warm', 'Cold', 'Warm', 'Warm'],
```

```
    'Humidity': ['Normal', 'High', 'High', 'High', 'Normal'],
```

```
    'Wind': ['Strong', 'Strong', 'Strong', 'Strong', 'Weak'],
```

```
    'Water': ['Warm', 'Warm', 'Warm', 'Warm', 'Cool'],
```

```
    'Forecast': ['Same', 'Same', 'Change', 'Same', 'Same'],
```

```
    'EnjoySport': ['yes', 'yes', 'no', 'yes', 'yes']
```

```
}
```

```
examples = pd.DataFrame(data)
```

```
# Apply Candidate-Elimination algorithm
```

```
S, G = candidate_elimination(examples)
```

```
print("Most Specific Hypothesis (S):", S)
```

```
print("Most General Hypothesis (G):", G)
```

```
#output
```

```
Most Specific Hypothesis (S): []
```

```
Most General Hypothesis (G): [('Sunny', '?', '?', '?', '?', '?'), ('?', 'Warm', '?', '?', '?', '?'), ('?', '?', '?', '?', '?', 'Same')]
```